

Hypervisor-Level Debugger with Radare2

r2vmi




Mathieu Tarral



Whoami



- Researcher at **F-Secure** 
- Malware behavioral analysis
- Maintainer of Nitro framework
- Created the KVM-VMI organization on Github

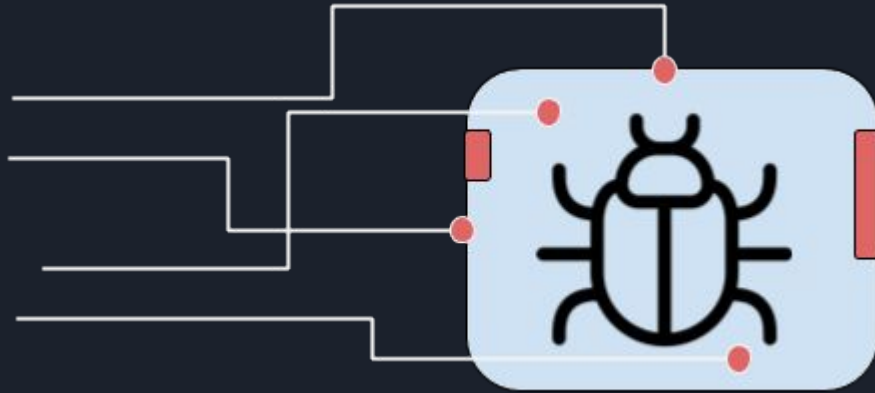
@mtaral 

Why ?



Debuggers are noisy

- A debugger modifies the execution environment of a debuggee



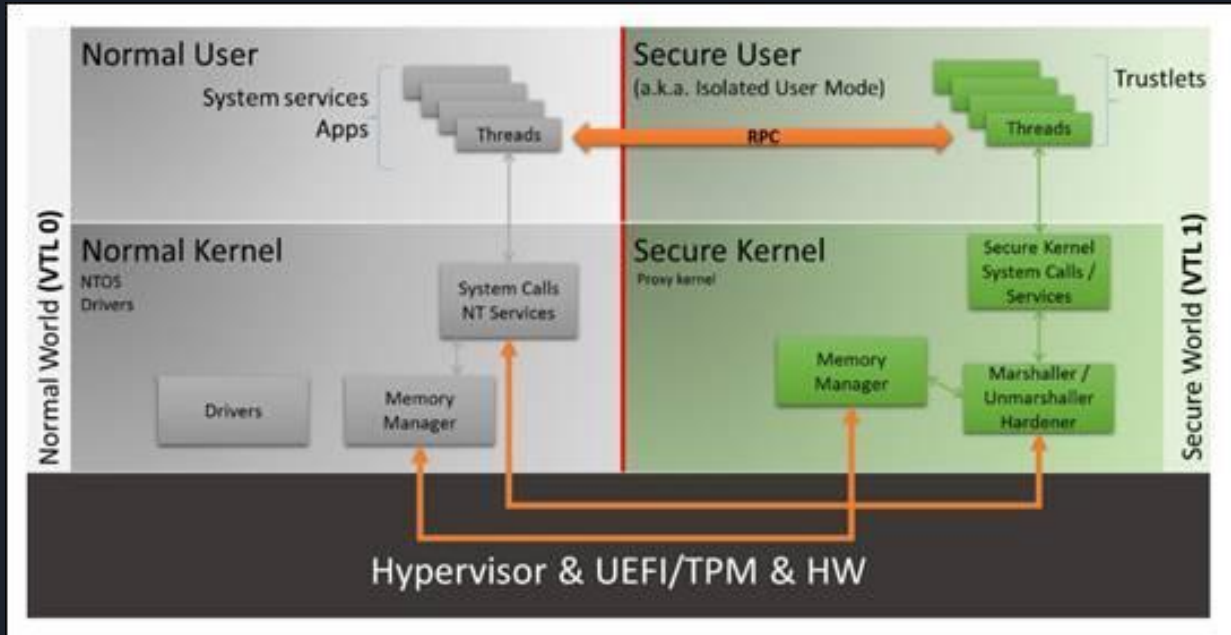


Debugger detection for dummies

- Windows API
 - `IsDebuggerPresent()`
 - `CheckRemoteDebuggerPresent()`
- Checking memory structures
 - `PEB.BeingDebugged`
 - `NtGlobalFlag`
 - `Process Heap Flags` and `ForceFlags`
 - scan for software breakpoints
- Side effects in system calls behavior
 - `NtClose` with invalid handle

Debugger's system view is incomplete

- Remote WinDBG is **not enough**





Moving to ring -1

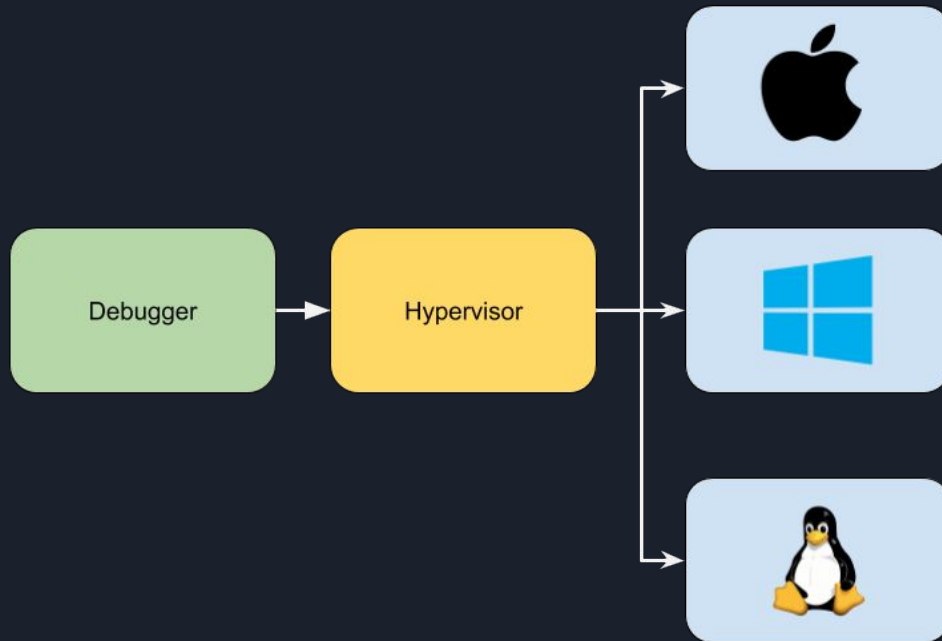
- **move** the debugging process to the hypervisor
- **Stealth**
 - **do not** use the operating system's debug API
 - **bonus**: invisible breakpoints with EPT violations
- **Full system analysis**
 - VMM's property: Resource control / Safety
 - access to the entire guest state
 - **bonus**: debug bootloaders



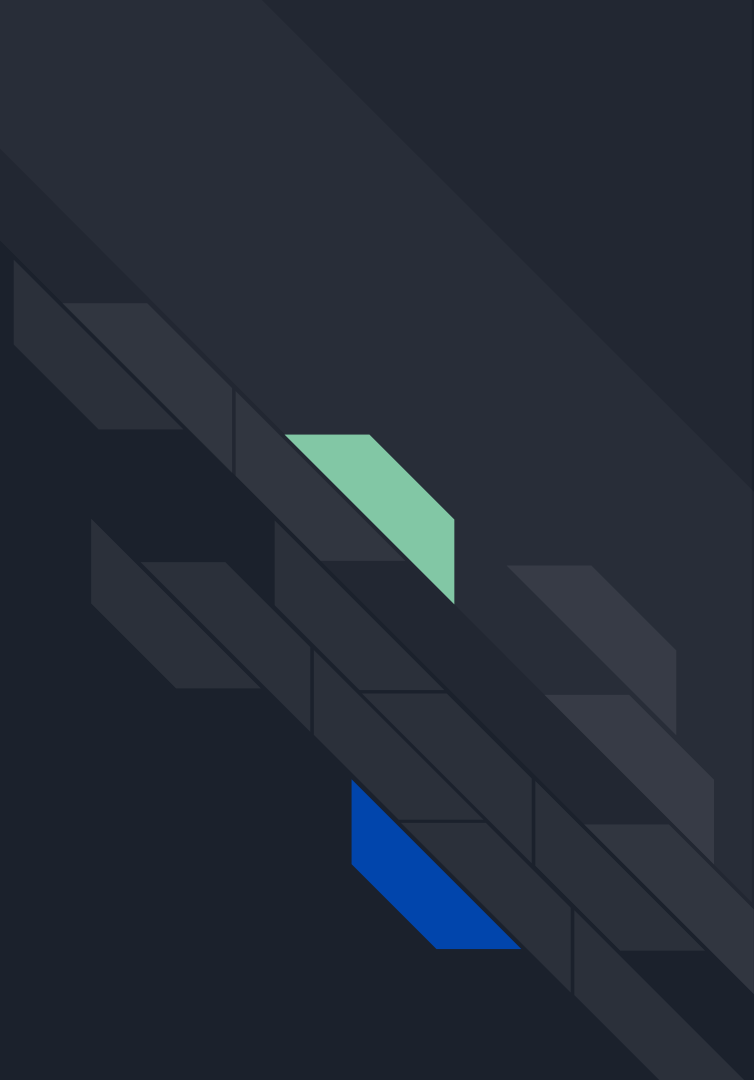
Zero configuration

- No remote debug agent/stub
- No custom VM setup
 - specific hardware
 - network card
 - serial cable
 - configuration
 - bcdedit /set debug on
 - bcdedit /dbgsettings serial debugport:1 baudrate:115200
- On-the-fly debugging

Cross-platform debugger



Projects ?





Hypervisor-based debuggers

- Debugger is built inside a tiny hypervisor
- HyperDBG (2010)
 - virtualize your host on-the-fly
- PulseDBG (2017)
 - load EFI bootloader (bootx64.efi)
 - client/server, needs 2 physical machines
- **Cons:**
 - debug real hardware
 - not designed to work with virtual machines



VMI-based debuggers

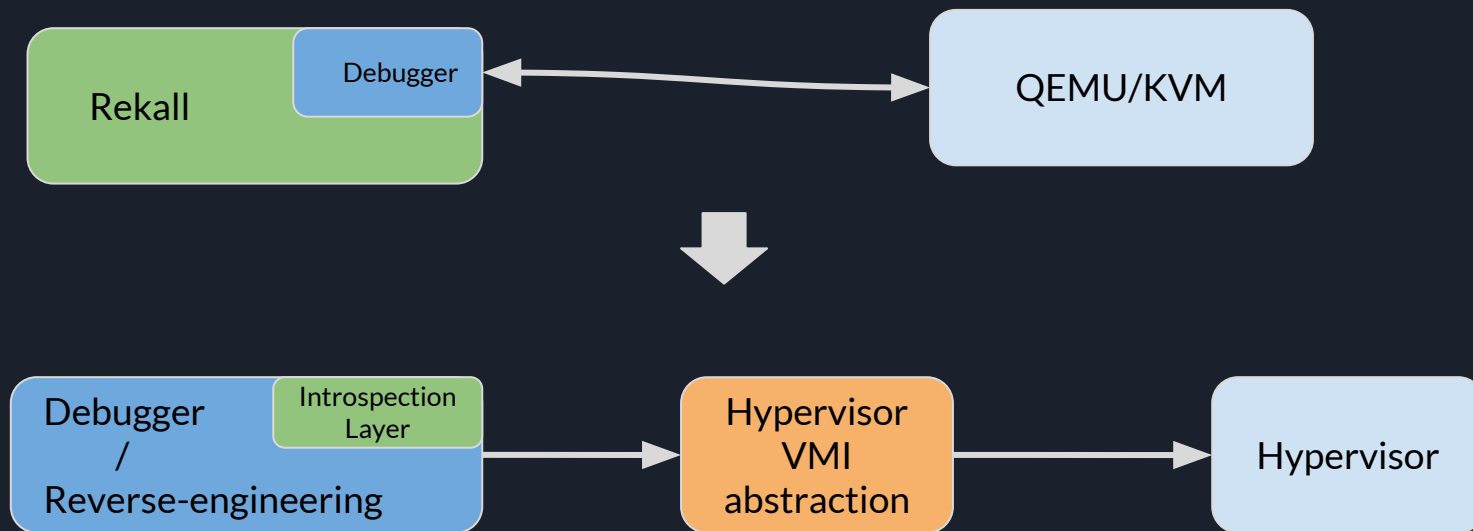
- Introspect the guest using VMI API's
- Talos PyREBox (2017)
 - dynamic instrumentation engine, based on QEMU
 - volatility for VMI
 - python callbacks
 - **cons:**
 - emulation
- FireEye rVMI (2017)
 - VMI on KVM
 - rekall as a debugger
 - **cons:**
 - KVM-only
 - maintainability?

How ?



Design

- Use VMI approach
- Improve rVMI





Hypervisor-agnostic: LibVMI

- VMI Abstraction layer
- Takes care of low-level details
- Offers basic introspection

	VCPU Registers	Physical memory	Hardware events
Xen	✓	✓	✓
KVM	✓	✓	✗

Architecture

- IO plugin (`io_vmi.c`)
 - initialize LibVMI, access memory and registers
- Debug plugin (`debug_vmi.c`)
 - attach process
 - singlestep
 - breakpoints
- `r2 -d vmi://vm_name:name|pid`



Implementation ?

The background features a series of dark gray, three-dimensional rectangular planes that recede into the distance, creating a sense of depth. A light green parallelogram is positioned on one of the upper planes, and a blue parallelogram is on a lower plane, both appearing to be part of the geometric structure.

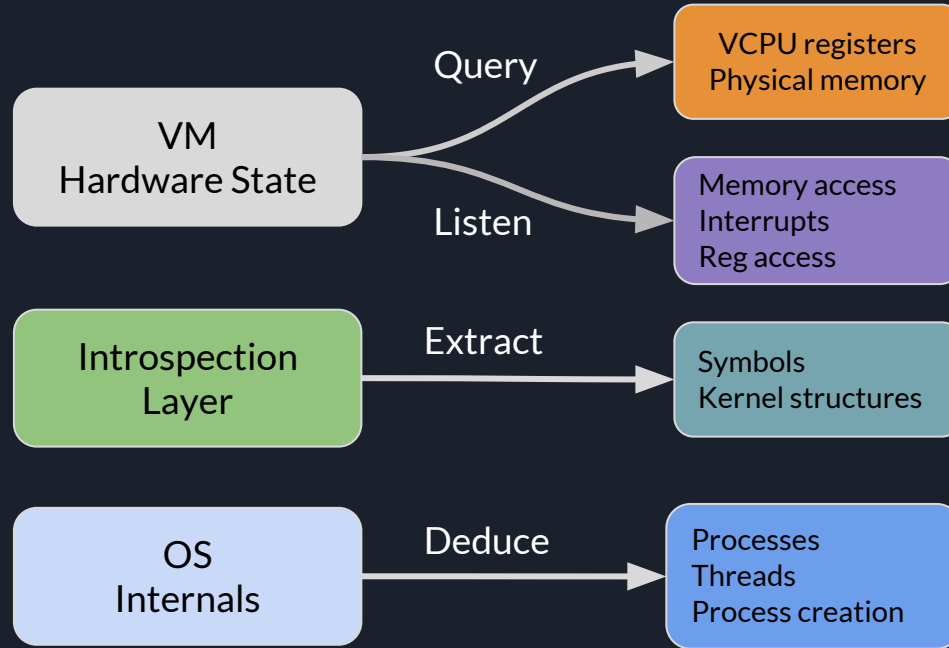


Challenges

- Debugging is a set of API provided by the operating system
- We have to rebuild these APIs from the hypervisor

```
//  
// Main User-Mode Debugger Loop //  
CreateProcess("target.exe", ..., DEBUG_PROCESS, ...);  
while (1)  
{  
    WaitForDebugEvent(&event, ...);  
    switch (event)  
    {  
        case ModuleLoad:  
            Handle/Ignore;  
            break;  
        case TerminateProcess:  
            Handle/Ignore;  
            break;  
        case Exception (code breakpoint, single step, etc...):  
            Handle/Ignore;  
            break;  
    }  
    ContinueDebugEvent(...);  
}
```

Challenges - What we have



Attaching existing process

- How to implement `_attach(RDebug *dbg, int pid) ?`
- Listen on MOV-TO-CR3 events
- Check for your targeted pid

```
[0xffffffff800026de1a9]> pd 10 @-3
0xffffffff800026de1a6      0f22da      mov cr3, rdx
;-- rip:
0xffffffff800026de1a9      4c8bbb88feff. mov r15, qword [rbx - 0x178]
0xffffffff800026de1b0      488b6e28     mov rbp, qword [rsi + 0x28]
0xffffffff800026de1b4      49896f04     mov qword [r15 + 4], rbp
0xffffffff800026de1b8      48896b28     mov qword [rbx + 0x28], rbp
0xffffffff800026de1bc      f7057e741800. test dword [0xffffffff80002865644], 4
,=< 0xffffffff800026de1c6      0f850b010000 jne 0xffffffff800026de2d7
0xffffffff800026de1cc      c6474900     mov byte [rdi + 0x49], 0
0xffffffff800026de1d0      0fbe86650100. movsx eax, byte [rsi + 0x165]
0xffffffff800026de1d7      84c0         test al, al
```



Introspection

- Use Rekall forensic framework
- Solution 1:
 - load a Python interpreter
 - call Rekall functions using Python C APIs (PyObject_CallObject)
- Solution 2:
 - run a Python RPC-endpoint, RESTful API, running Rekall
 - serialize objects and structs, send JSON

Introspection

- Solution 3:

- Load kernel symbols from LibVMI (r2>. \symbols)
- run a Rekall interactive session using VMIAddressSpace
- `rekall -f vmi://xen/vm_name`

```
[1] xenwin7 21:38:31> pslist
-----> pslist()2800000 (0xa000000)
 _EPROCESS      name          pid  ppid  thread_count handle_count session_id wow64  process_create_time
-----
```

0xfa80033b3040	System	4	0	76	456	-	False	2018-08-22 01:18:51Z
0xfa80045b0b30	svchost.exe	192	464	17	312	0	False	2018-08-22 01:18:58Z
0xfa80044ef8a0	smss.exe	252	4	2	29	-	False	2018-08-22 01:18:51Z
0xfa80042fd060	spoolsv.exe	296	464	12	264	0	False	2018-08-22 01:18:58Z
0xfa800469e420	csrss.exe	324	316	8	278	0	False	2018-08-22 01:18:53Z
0xfa800459e810	svchost.exe	340	464	16	401	0	False	2018-08-22 01:18:57Z
0xfa800470f060	wininit.exe	372	316	3	74	0	False	2018-08-22 01:18:54Z
0xfa8004839060	csrss.exe	380	364	7	136	1	False	2018-08-22 01:18:54Z
0xfa8004852460	winlogon.exe	420	364	3	107	1	False	2018-08-22 01:18:54Z
0xfa800488ab30	services.exe	464	372	5	178	0	False	2018-08-22 01:18:54Z
0xfa8004892b30	lsass.exe	472	372	6	527	0	False	2018-08-22 01:18:54Z
0xfa800465eac0	lsn.exe	480	372	9	135	0	False	2018-08-22 01:18:54Z



Attaching new process

- How to implement `_attach(RDebug *dbg, int pid) ?`
 - Find CR3 doesn't exist in kernel process list **yet**
- How to access the entrypoint ?
- Windows Internals

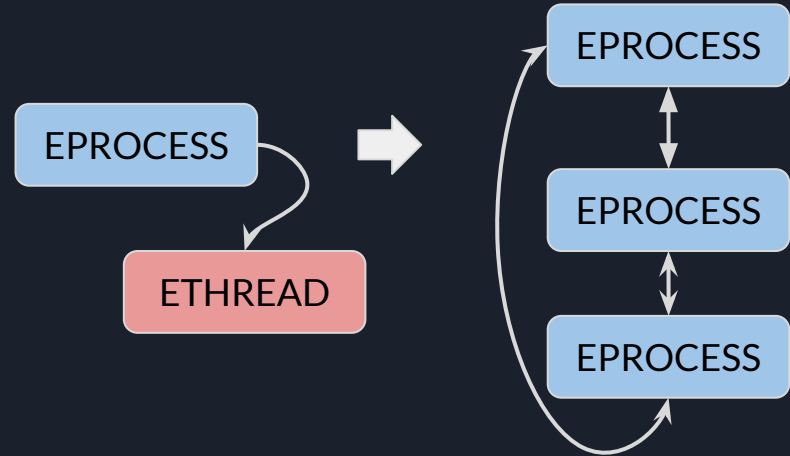
Process Creation (Windows 7)

- `nt!NtCreateUserProcess`

- `pspAllocateProcess`
- `pspsInsertProcess`
- `PspAllocateThread`
- `PspInsertThread`

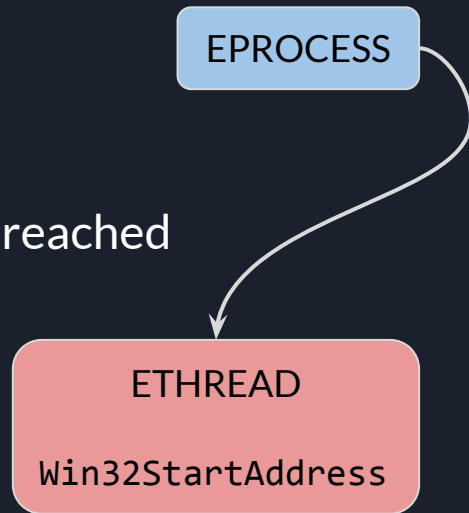
- `nt!KiStartUserThread`

- `ntdll!LdrInitializeThunk`
- `ntdll!NtContinue`
- `ntdll!RtlUserThreadStart`
 - `call entrypoint!`



To the entrypoint

- Read `Win32StartAddress`
- Populated during `nt!pspAllocateThread`
- Available when `nt!KiStartUserThread` is reached
- Solution:
 - break on `nt!KiStartUserThread`
 - read `Win32StartAddress`
 - break on entrypoint
 - **not** mapped yet !
 - Tests for the page to be mapped
 - `nt!MmAccessFault`
 - `singlestep` until `ring3`
 - watch for page tables modification



Status ?





Features

- Intercept a existing process by Name/pid
- Single-step process execution
- Set memory breakpoints
- Load Rekall symbols (kernel only)
- Radare2 interface !



Issues

- Attach existing process
 - Find threads context, find rip
- Attach new process
 - debug VMI state of guest
 - Xen development
- Break on addresses not mapped yet
 - Watch page tables for modification
- Introspection
 - Access Rekall from C ?
 - Python Rabin plugin + Rekall VMIAddressSpace ?

Demo



Future ?





Future

- Malware analysis
 - stealth sandboxes
 - highly interactive reverse-engineering frameworks
- Vulnerability research
- Multi-purpose, cross-plaform, full system debugger



<https://github.com/Wenzel/r2vmi>





Thanks

- pancake (*radare2*)
- Maxime Morin (*radare2*)
- Tamas K Lengyel (*LibVMI*)
- Michael Cohen (*Rekall*)
- Alexey Konovalov (*Windows Internals*)

Questions

