

EVM Bytecode and smart contract internals

Fedor Sakharov
fedor.sakharov@gmail.com



Agenda

- About me
- A bit about Ethereum
- Smart contracts and their bytecode
- Implementing EVM support in r2
- Future work



About me

- @m0nt3kk1 on twitter, @montekki on github



About me

- @m0nt3kk1 on twitter, @montekki on github
- Code mostly in C or Go, sys and net stuff



About me

- @m0nt3kk1 on twitter, @montekki on github
- Code mostly in C or Go, sys and net stuff
- @xvilka brought me to r2 ~5 years ago



About me

- @m0nt3kk1 on twitter, @montekki on github
- Code mostly in C or Go, sys and net stuff
- @xvilka brought me to r2 ~5 years ago
- Made a living working in infosec and crypto



About me

- @m0nt3kk1 on twitter, @montekki on github
- Code mostly in C or Go, sys and net stuff
- @xvilka brought me to r2 ~5 years ago
- Made a living working in infosec and crypto
- Now unemployed



A bit about Ethereum



Ethereum: known facts

- Distributed blockchain-based computing platform and operating system featuring smart contract functionality (wiki)



Ethereum: known facts

- Distributed blockchain-based computing platform and operating system featuring smart contract functionality (wiki)
- Often referred to as blockchain 2.0 (hype)



Ethereum: known facts

- Distributed blockchain-based computing platform and operating system featuring smart contract functionality (wiki)
- Often referred to as blockchain 2.0 (hype)
- p2p blockchain stuff + Turing-complete VM (reality)



Ethereum: known facts

- Distributed blockchain-based computing platform and operating system featuring smart contract functionality (wiki)
- Often referred to as blockchain 2.0 (hype)
- p2p blockchain stuff + Turing-complete VM (reality)
- Some money inside ~\$30 000 000 000



Ethereum: known facts

- Distributed blockchain-based computing platform and operating system featuring smart contract functionality (wiki)
- Often referred to as blockchain 2.0 (hype)
- p2p blockchain stuff + Turing-complete VM (reality)
- Some money inside ~\$30 000 000 000



Ethereum: known hacks

- Parity Freeze ~400 000 000 \$ freezed
- Parity Multisig hack ~32 000 000 \$ stolen
- DAO hack ~70 000 000 \$ stolen



Smart contracts and their bytecode



Smart contracts

- Smart contracts are a digitized version of real-world contracts



Smart contracts

- Smart contracts are a digitized version of real-world contracts
- In Eth different SC languages exist



Smart contracts

- Smart contracts are a digitized version of real-world contracts
- In Eth different SC languages exist
- Every one of them is translated into EVM bytecode



Smart contracts

- Runtime environment is the Ethereum VM



Smart contracts

- Runtime environment is the Ethereum VM
- Specified in the Yellow Paper



Smart contracts

- Runtime environment is the Ethereum VM
- Specified in the Yellow Paper
- A stack-based machine



Smart contracts

- Runtime environment is the Ethereum VM
- Specified in the Yellow Paper
- A stack-based machine
- Code always runs as intended (no side effects)



Smart contracts

- Runtime environment is the Ethereum VM
- Specified in the Yellow Paper
- A stack-based machine
- Code always runs as intended (no side effects)
- Code always terminates



Smart contracts: example

```
pragma solidity ^0.4.0;

contract Example {
    uint a;

    function setA(uint b) public {
        a = b;
    }
}
```

- Solidity code has no single entry point
- Transactions to contract call it's methods with some arguments
- Code is compiled to bytecode which is stored on-chain with it's abi



Smart contracts: example

```
pragma solidity ^0.4.0;
```

```
contract Example {
```

```
    uint a;
```

```
    function setA(uint b) public {
```

```
        a = b;
```

```
    }
```

```
}
```

```
$ solc ./example.sol --abi -o ./out/
```

```
$ cat ./out/Example.abi | pjson
```

```
[{  
  "constant": false,  
  "inputs": [  
    {  
      "name": "b",  
      "type": "uint256"  
    }  
  ],  
  "name": "setA",  
  "outputs": [],  
  "payable": false,  
  "stateMutability": "nonpayable",  
  "type": "function"  
}]
```



Smart contracts: example

```
pragma solidity ^0.4.0;

contract Example {

    uint a;

    function setA(uint b) public {

        a = b;

    }

}
```

```
$ solc ./example.sol --bin-runtime -o
./out/
```

```
$ cat ./out/Example.bin-runtime
```

```
608060405260043610603f576000357c0100000
0000000000000000000000000000000000000
000000000000900463ffffffff168063ee919d5
0146044575b600080fd5b348015604f57600080
fd5b50606c60048036038101908080359060200
190929190505050606e565b005b806000819055
50505600a165627a7a72305820c4b0ba922b4d
8b5902a3090f2e7e045ffcfe0c553ebe553c
decabfca42c50029
```



Smart contracts: example

- Hex-encoded bytecode of the Smart Contract
- Execution starts from the first byte
- Input is stored in the transaction parameters

```
$ solc ./example.sol --bin-runtime -o  
./out/
```

```
$ cat ./out/Example.bin-runtime
```

```
608060405260043610603f576000357c0100000  
000000000000000000000000000000000000  
000000000000900463ffffffff168063ee919d5  
0146044575b600080fd5b348015604f57600080  
fd5b50606c60048036038101908080359060200  
190929190505050606e565b005b806000819055  
50505600a165627a7a72305820c4b0ba922b4d  
8b5902a3090f2e7e045ffcfe0c553ebe553c  
decabfca42c50029
```



Smart contracts: run & develop

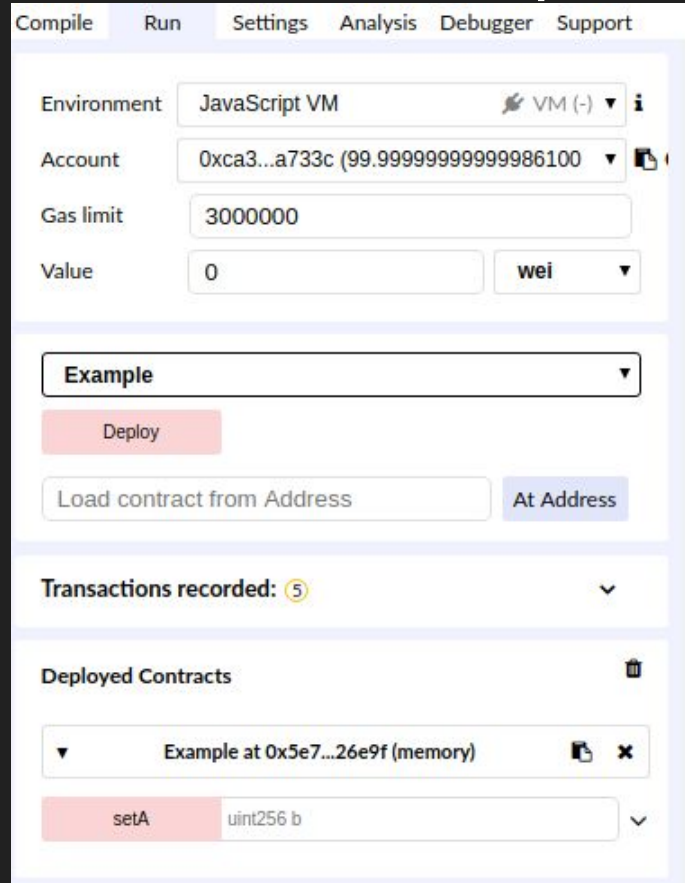
The widely-used IDE for Solidity is Remix
(<https://remix.ethereum.org>)

- Editor
- Compiler
- Debugger
- Embedded VM



Smart contracts: run and develop

You have to deploy
the contract and
then you can call
it's methods





The screenshot displays the Remix IDE interface with the following components:













- Navigation Bar:** Compile, Run, Settings, Analysis, Debugger, Support.
- Environment:** JavaScript VM (with a VM icon and a dropdown arrow).
- Account:** 0xca3...a733c (99.99999999999986100) (with a dropdown arrow and a document icon).
- Gas limit:** 3000000 (in a text input field).
- Value:** 0 (in a text input field) and wei (in a dropdown menu).
- Example:** A dropdown menu showing 'Example'.
- Deploy:** A red button.
- Load contract from Address:** A text input field.
- At Address:** A blue button.
- Transactions recorded:** 5 (with a dropdown arrow).
- Deployed Contracts:** A section with a trash icon.
 - Example at 0x5e7...26e9f (memory):** A dropdown menu with a document icon and a close icon (X).
 - setA:** A red button.
 - uint256 b:** A text input field with a dropdown arrow.



Smart contracts: run & develop


Calling setA method displays TX summarized:

 [vm] from:0xca3...a733c to:Example.setA(uint256) 0x5e7...26e9f value:0 wei data:0xee9...0002a logs:0 hash:0x794...ac2fd Debug 

status	0x1 Transaction mined and execution succeed
transaction hash	0x7945500403636dffa934de05968b50221b79230587c01683d8d44f16bf9ac2fd 
from	0xca35b7d915458ef540ade6068dfe2f44e8fa733c 
to	Example.setA(uint256) 0x5e72914535f202659083db3a02c984188fa26e9f 
gas	3000000 gas 
transaction cost	41675 gas 
execution cost	20211 gas 
hash	0x7945500403636dffa934de05968b50221b79230587c01683d8d44f16bf9ac2fd 
input	0xee9...0002a 
decoded input	<pre>{ "uint256 b": "42" }</pre> 
decoded output	<pre>{}</pre>
logs	<pre>[]</pre>  
value	0 wei 



instance:



EVM Bytecode



EVM bytecode

- Evm is a stack-based machine
- During runtime data is stored in memory or stack
- Data is stored permanently in contract's storage on-chain

PUSH1 32	Stack: [32]
PUSH1 42	Stack: [42, 32]
ADD	Stack: [74]



EVM bytecode

- Flow control is done pretty conventionally
- Beginnings of the basic blocks are JUMPDEST instructions

PUSH1 01	Stack: [1], PC: 0
JUMPDEST	Stack: [1], PC: 2
PUSH1 02	Stack: [1, 2], PC: 3
ADD	Stack: [3], PC: 5
PUSH1 0x2	Stack: [3, 2], PC: 6
JUMP	Stack: [3], PC: 7
PUSH1 02	Stack: [3, 2], PC: 3
ADD	Stack: [5], PC: 5



EVM bytecode

Two types of instructions:

1. Ones you expect to find in a stack-based VM
2. Ethereum-specific instructions:

CREATE, CALL, CALLCODE, DELEGATECALL, REVERT,
SUICIDE, etc.



EVM bytecode

The code in the contract is always executed from the 0x0000 addr

First instructions load TX data into VM memory and do other initialization

```
(fcn) fcn.00000000 121
      fcn.00000000 ();
      0x00000000      6080      push1 0x80
      0x00000002      6040      push1 0x40
      0x00000004      52        mstore
      0x00000005      6004      push1 0x4
      0x00000007      36        calldatasize
      0x00000008      10        lt
      0x00000009      603f      push1 0x3f
      < 0x0000000b      57        jumpi
      0x0000000c      6000      push1 0x0
      0x0000000e      35        calldataload
      0x0000000f      7c0100000000. push29 0x0
      0x0000002d      90        swap1
      0x0000002e      04        div
      0x0000002f      63ffffffff      push4 0xffffffff
      0x00000034      16        and
```



EVM bytecode

This code will load the data into memory

And it will AND this data with 0xffffffff thus taking the first four bytes of it

(fcn) fcn.00000000 121		
fcn.00000000 ();		
0x00000000	6080	push1 0x80
0x00000002	6040	push1 0x40
0x00000004	52	mstore
0x00000005	6004	push1 0x4
0x00000007	36	calldatasize
0x00000008	10	lt
0x00000009	603f	push1 0x3f
0x0000000b	57	jumpi
0x0000000c	6000	push1 0x0
0x0000000e	35	calldataload
0x0000000f	7c0100000000.	push29 0x0
0x0000002d	90	swap1
0x0000002e	04	div
0x0000002f	63ffffffff	push4 0xffffffff
0x00000034	16	and



EVM bytecode

Those value is actually a hash signature of the function

A hash signature is the first four bytes of keccak hash of the function signature

	0x00000035	80	dup1
	0x00000036	63ee919d50	push4 0xee919d50
	0x0000003b	14	eq
	0x0000003c	6044	push1 0x44
	0x0000003e	57	jumpi
<	0x0000003f	5b	jumpdest
	0x00000040	6000	push1 0x0
	0x00000042	80	dup1
	0x00000043	fd	revert
	0x00000044	5b	jumpdest
	0x00000045	34	callvalue

```
keccak256("setA(uint256)")
```

```
ee919d50445cd9f463621849366a537968fe1ce096894b0d0c001528383d4769
```



EVM bytecode

Those value is actually a hash signature of the function

A hash signature is the first four bytes of keccak hash of the function signature

0x00000035	80	dup1
0x00000036	63ee919d50	push4 0xee919d50
0x0000003b	14	eq
0x0000003c	6044	push1 0x44
< 0x0000003e	57	jumpi
└─> 0x0000003f	5b	jumpdest
0x00000040	6000	push1 0x0
0x00000042	80	dup1
0x00000043	fd	revert
└─> 0x00000044	5b	jumpdest
0x00000045	34	callvalue

keccak256("setA(uint256)")

ee919d50445cd9f463621849366a537968fe1ce096894b0d0c001528383d4769



EVM bytecode

If the first four bytes of the input data doesn't contain a hash known to the function dispatcher, the execution is terminated and the transaction is reverted

0x00000035	80	dup1
0x00000036	63ee919d50	push4 0xee919d50
0x0000003b	14	eq
0x0000003c	6044	push1 0x44
< 0x0000003e	57	jumpi
└─> 0x0000003f	5b	jumpdest
0x00000040	6000	push1 0x0
0x00000042	80	dup1
0x00000043	fd	revert
└─> 0x00000044	5b	jumpdest
0x00000045	34	callvalue



EVM bytecode

And after some other checks and init code we end up in our function

it adds 0x42 to an input parameter

and stores it if you recall

```
function setA(uint b) public {  
    a = b + 0x42;  
}
```


└─> 0x0000006e	5b	jumpdest
0x0000006f	6042	push1 0x42
0x00000071	81	dup2
0x00000072	01	add
0x00000073	6000	push1 0x0
0x00000075	81	dup2
0x00000076	90	swap1
0x00000077	55	sstore
0x00000078	50	pop
0x00000079	50	pop
0x0000007a	56	jump
0x0000007b	00	stop



EVM bytecode: specification experience

60s & 70s: Push Operations

Value	Mnemonic	δ	α	Description
0x60	PUSH1	0	1	<p>Place 1 byte item on stack.</p> $\mu'_s[0] \equiv c(\mu_{pc} + 1)$ <p>where $c(x) \equiv \begin{cases} I_b[x] & \text{if } x < \ I_b\ \\ 0 & \text{otherwise} \end{cases}$</p> <p>The bytes are read in line from the program code's bytes array. The function c ensures the bytes default to zero if they extend past the limits. The byte is right-aligned (takes the lowest significant place in big endian).</p>
0x61	PUSH2	0	1	<p>Place 2-byte item on stack.</p> $\mu'_s[0] \equiv c((\mu_{pc} + 1) \dots (\mu_{pc} + 2))$ <p>with $c(x) \equiv (c(x_0), \dots, c(x_{\ x\ -1}))$ with c as defined as above. The bytes are right-aligned (takes the lowest significant place in big endian).</p>
\vdots	\vdots	\vdots	\vdots	\vdots



Implementing EVM support in R2



ASM plugin

- The opcodes and the commands format is straightforward
- Opcodes are 1-byte
- Most of them have no operands
- Just follow the yellowpaper



ANAL plugin

- Things become more tricky
- All jumps take values from stack as operands
- In general case not possible to retrieve the dst addr of the jump during the analysis
- If previous command is PUSH_*, take dst addr from it



ASM + ANAL: let's take a look

```
theodor@fran:r2presentation$ rax2 -s < ./out/Example.bin-runtime > ./out/Example.bin-runtime.bin
theodor@fran:r2presentation$ r2 -a evm ./out/Example.bin-runtime.bin
-- Trace register changes while debugging with 'e trace.cmtregs=true'
[0x00000000]> aa
[×] Analyze all flags starting with sym. and entry0 (aa)
[0x00000000]> pd 20
(fcn) fcn.00000000 118
    fcn.00000000 ();
    3x00000000      6080      push1 0x80
    0x00000002      6040      push1 0x40
    0x00000004      52       mstore
    0x00000005      6004      push1 0x4
    0x00000007      36       calldatasize
    0x00000008      10       lt
    0x00000009      603f      push1 0x3f
    0x0000000b      57       jumpi
    0x0000000c      6000      push1 0x0
    0x0000000e      35       calldataload
    0x0000000f      7c0100000000. push29 0x0
    0x0000002d      90       swap1
    0x0000002e      04       div
    0x0000002f      63ffffffff      push4 0xffffffff
    0x00000034      16       and
    0x00000035      80       dup1
    0x00000036      63ee919d50      push4 0xee919d50
    0x0000003b      14       eq
    0x0000003c      6044      push1 0x44
    0x0000003e      57       jumpi
[0x00000000]> █
```



ASM + ANAL: let's take a look

```
theodor@fran:r2presentation$ rax2 -s < ./out/Example.bin-runtime > ./out/Example.bin-runtime.bin
theodor@fran:r2presentation$ r2 -a evm ./out/Example.bin-runtime.bin
-- Trace register changes while debugging with 'e trace.cmtregs=true'
```

```
[0x00000000]> aa
```

```
[x] Analyze all flags starting with sym. and entry0 (aa)
```

```
[0x00000000]> pd 20
```

```
(fcn) fcn.00000000 118
    fcn.00000000 ();
    3x00000000      6080      push1 0x80
    0x00000002      6040      push1 0x40
    0x00000004      52       mstore
    0x00000005      6004      push1 0x4
    0x00000007      36       calldatasize
    0x00000008      10       lt
    0x00000009      603f      push1 0x3f
    0x0000000b      57       jumpi
    0x0000000c      6000      push1 0x0
    0x0000000e      35       calldataload
    0x0000000f      7c0100000000. push29 0x0
    0x0000002d      90       swap1
    0x0000002e      04       div
    0x0000002f      63ffffffff      push4 0xffffffff
    0x00000034      16       and
    0x00000035      80       dup1
    0x00000036      63ee919d50      push4 0xee919d50
    0x0000003b      14       eq
    0x0000003c      6044      push1 0x44
    0x0000003e      57       jumpi
```

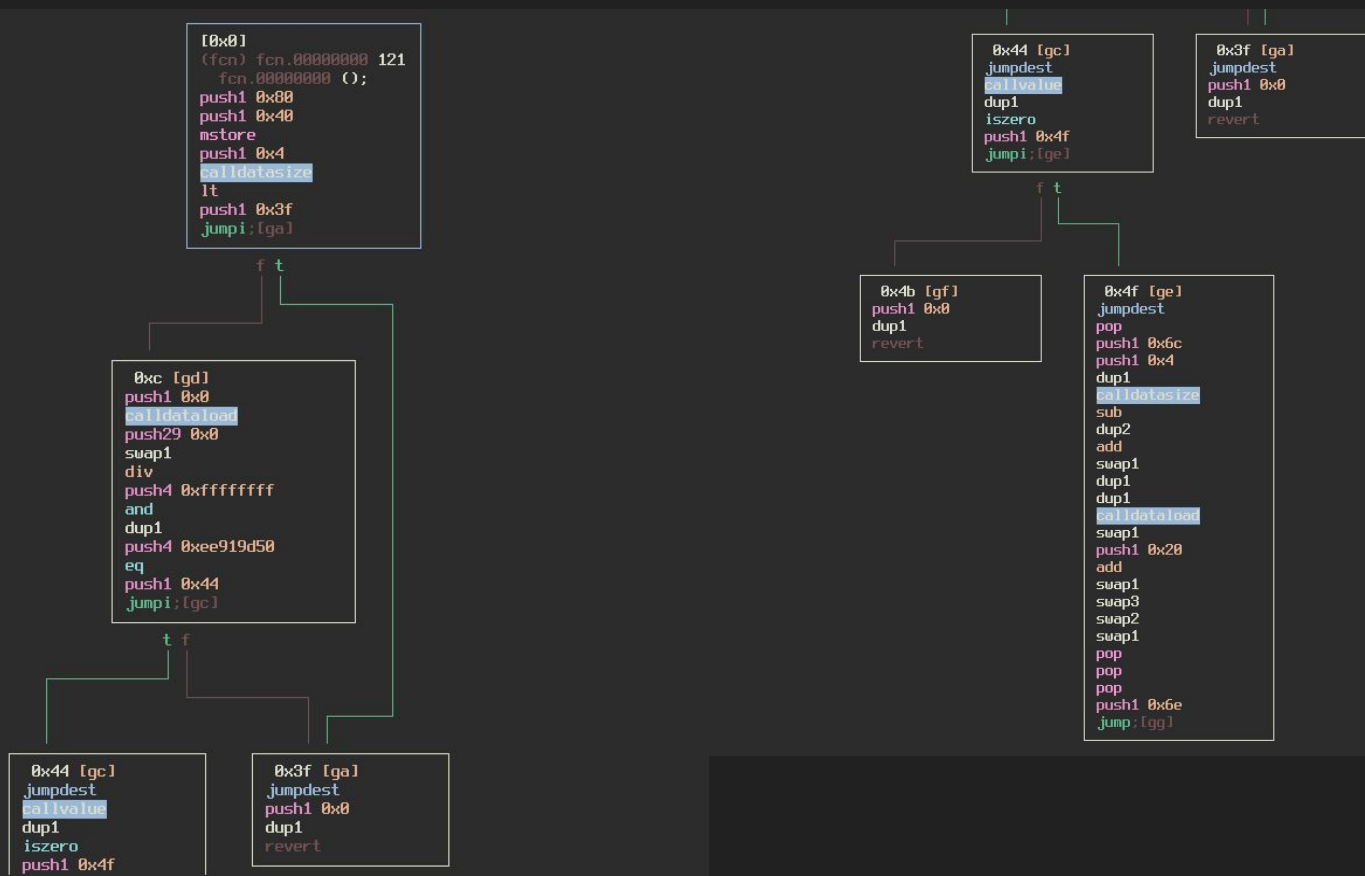
```
[0x00000000]> █
```

Well, at least it's usable

However, the long-operand commands like push29 are not displayed correctly



ASM + ANAL: let's take a look



IO plugin

- Contract lives on-chain, why not read it from there?



IO plugin

- Contract lives on-chain, why not read it from there?
- Go-ethereum exposes a set of RPC methods



IO plugin

- Contract lives on-chain, why not read it from there?
- Go-ethereum exposes a set of RPC methods, `eth_GetCode`, for instance

```
curl -X POST --data
'{"jsonrpc": "2.0", "method": "eth_getCode", "params": ["0xa94f5374f5e5edbc8e2a8697c15331677e6ebf0b",
"0x2"], "id": 1}'

// Result
{
  "id": 1,
  "jsonrpc": "2.0",
  "result":
"0x600160008035811a818181146012578301005b601b6001356025565b8060005260206000f25b600060078202905091905056"
}
```



IO plugin

- Contract lives on-chain, why not read it from there?
- Go-ethereum exposes a set of RPC methods, `eth_GetCode`, for instance
- Bring in `libcurl` and a json parser to the plugin



IO plugin

- Contract lives on-chain, why not read it from there?
- Go-ethereum exposes a set of RPC methods, `eth_GetCode`, for instance
- Bring in `libcurl` and a json parser to the plugin
- Now we can talk to the node



I0 plugin

```
$ r2 -a evm "evm://localhost:8545@0x075121e8f930cb7bc21cc726600e532b3e60e7d1"
= attach 1 1
-- In Soviet Russia, radare2 has documentation.
[0x00000000]> aa
[x] Analyze all flags starting with sym. and entry0 (aa)
[0x00000000]> pd 10
    fcn.eax ();
        0x00000000      6080      push1 0x80
        0x00000002      6040      push1 0x40
        0x00000004      52        mstore
        0x00000005      6004      push1 0x4
        0x00000007      36        calldatasize
        0x00000008      10        lt
        0x00000009      603f      push1 0x3f
        0x0000000b      57        jumpi
        0x0000000c      6000      push1 0x0
        0x0000000e      35        calldataload
```



DBG plugin

- EVM has no debugging std debugging interfaces



DBG plugin

- EVM has no debugging std debugging interfaces
- You can not trace the live execution of a call to smart contract



DBG plugin

- EVM has no debugging std debugging interfaces
- One can not trace the live execution of a call to smart contract
- However, RPC API allows to get a trace of any transaction

```
{"method": "debug_traceTransaction", "params": [txHash, {}]}
```



DBG Plugin

```
{"method": "debug_traceTransaction", "params": [txHash, {}]}
```

Will return a step-by-step trace of a transaction in JSON form

```
...  
{  
  gas: 85301,  
  returnValue: "",  
  structLogs: [{  
    depth: 1,  
    error: "",  
    gas: 162106,  
    gasCost: 3,  
    memory: null,  
    op: "PUSH1",  
    pc: 0,  
    stack: [],  
    storage: {}  
  }],  
  ...  
}
```



DBG Plugin

- Read and parse the whole transaction trace
- Implement DBG plugin API
- Emulate “live” execution, actually seeking forward in the trace
- Breakpoints are just checking if the current element in the trace equals bpt addr
- And so on



Future work



Future work

- ESIL looks promising, since we can emulate the finite executions well
- R2dec support also, the code looks simple to decompile
- Support for displaying 32-byte operands
- Better JMP/CALL/RET dst computation



Useful links

- https://medium.com/@theo_montekki my blog with written info
- [A series of blog posts on blog.zepplin.solutions](https://blog.zepplin.solutions)
- **CTF** <https://ethernaut.zepplin.solutions/>
- <https://github.com/ethereum/wiki/wiki/Ethereum-Contract-ABI>
- <http://yellowpaper.io/>
- <https://hackernoon.com/smart-contract-security-part-1-reentrancy-attacks-ddb3b2429302>
- <https://applicature.com/blog/history-of-ethereum-security-vulnerabilities-hacks-and-their-fixes>



Thank you!

Questions?

